

Modern C++ (C++11/14)

Geeks Anonymes


Damien Gerard

October 19th, 2016

Introduction

“C++11 feels like a new language.” Bjarne Stroustrup (inventor of C++)

Labels



since C++11



since C++14



since C++17



that ain't for kids

11

```
int myArray[5] = {1, 2, 3, 4, 5};
for(int x : myArray) {
    std::cout << x << std::endl;
}
for(int& x : myArray) {
    x *= 2; // double value and save
}
for(auto& x : myArray) { // auto = int
    x *= 2;
}
```

std.

11

`f(nullptr_t)` will actually call `f(int*)` using implicit conversion.

Useful albeit limited to Plain Old Data (POD) structs and classes.

11

```
std::vector<std::string> v = { "Hello", "world" };
std::vector<std::string> v({ "Hello", "world" });
std::vector<std::string> v{ "Hello", "world" };
```

```
int val = 5.2; // automatic narrowing
int val{5.2}; // error or warning: type 'double'
              // cannot be narrowed to 'int' in initializer list
              // insert an explicit static cast static_cast<int>( )
```

```
Widget getWidget() {
    return {0.43f, 10}; // no need for explicit type
}
```

Most vexing parse

```
struct Widget {  
    Widget();  
};  
struct WidgetKeeper {  
    WidgetKeeper(Widget w);  
};  
WidgetKeeper wKeeper(Widget()); // most vexing parse
```

Ambiguous call:

- 1 new instance of Widget, sent to constructor of WidgetKeeper;
- 2 function declaration: name "wKeeper", return type WidgetKeeper, single (unnamed) parameter which is a function returning type Widget.

Most vexing parse

Most developers expect the first, but the C++ standard requires it to be interpreted as the second.

Workaround in C++03 to ensure the first interpretation:

```
WidgetKeeper time_keeper( (Widget()) );
```

C++11's uniform initialization syntax:

```
WidgetKeeper time_keeper{Widget{}};
```

Returning multiple values

```
void divide(int dividend, int divisor, int& quotient,
           int& remainder)
{
    quotient = dividend / divisor;
    remainder = dividend % divisor;
}
```

Or even worse, return quotient through the returned value and remainder as a reference variable:

```
int divide(int dividend, int divisor, int& remainder) {
    remainder = dividend % divisor;
    return dividend / divisor;
}
```

Returning multiple values

Or, but rather heavy:

```
struct divideResult {  
    int quotient;  
    int remainder;  
};
```

```
divideResult divide(int dividend, int divisor) {  
    return { dividend / divisor, dividend % divisor };  
}
```

11

```
divideResult result = divide(10, 3);  
std::cout << result.quotient << std::endl;  
std::cout << result.remainder << std::endl;
```

Returning multiple values

```
std::tuple<int,int> divide(int dividend, int divisor) {  
    return std::make_tuple(dividend / divisor,  
                           dividend % divisor);  
}
```

11

```
int quotient, remainder;  
std::tie(quotient, remainder) = divide(10, 3);
```

C++17's structured bindings (Clang-4.0 only, as of today):

```
auto [quotient, remainder] = divide(10, 3);  
std::cout << quotient << std::endl;  
std::cout << remainder << std::endl;
```

17

std::unordered_*

11

std::set, map, multiset, multimap associative containers are implemented as **balanced trees**. The std::unordered_set, unordered_map, unordered_multiset, unordered_multimap alternatives are implemented as **hash tables**.

Including hash tables was one of the most recurring requests. It was not adopted in C++03 due to time constraints only. Although hash tables are less efficient than a balanced tree in the worst case (in the presence of many collisions), they perform better in many real applications.

emplace, emplace_back

11

emplace and emplace_back can construct an element *in-place*.

```
auto employees = std::unordered_map<int, std::string>{};
auto e1 = std::pair<int, std::string>{1, "John Smith"};

employees.insert(e1);
employees.insert(std::make_pair(2, "Mary Jones"));
employees.emplace(3, "James Brown"); // construct in-place

for (const auto& e : employees) {
    std::cout << e.first << ": " << e.second << std::endl;
}
```

In-class initialization

In C++03, in-class initialization on static const members of integral or enumeration type only.

```
struct Widget {  
    static const int a = 7;           // compiles  
    float b = 7.2;                   // error: not static const integral  
    const int c = 7;                 // error: not static  
    static int d = 7;                // error: not const  
    static const float e = 7.2;      // error: not integral  
    const static int arr[] = { 1, 2, 3 };  
    // error: must be initialized out of line  
};
```

In-class initialization

11

C++11 allows some of them:

```
struct Widget {  
    static const int a = 7;  
    float b = 7.2;  
    const int c = 7;  
    static int d = 7;      // still fails  
    static float e = 7.2; // still fails  
    constexpr static int arr[] = { 1, 2, 3 };  
    constexpr static std::complex<double> f = { 1, 2 };  
};
```

A non-const static variable still has to be initialized outside the class with

```
int Widget::d = 7;  
float Widget::e = 7.2;
```

ODR-use

R

An object is odr-used if its address is taken, or a reference is bound to it. If a const or constexpr static data member is odr-used, a redefinition (no initializer permitted) at namespace scope is required.

```
struct Widget {  
    static const int n = 1;  
    static constexpr int m = 4;  
};  
const int& f(const int& r);  
// call f(Widget::n) or f(Widget::m) somewhere  
// redefinition at namespace scope required:  
const int Widget::n, Widget::m;
```

11

C++17 alleviates this constraint for constexpr. A static data member declared constexpr is implicitly inline and needs not be redeclared at namespace scope.

17

constexpr

Declare something that can be evaluated down to a constant:

```
constexpr double pi() { return std::atan(1) * 4; }  
constexpr double getCircleSurface(double r) {  
    return pi() * r * r;  
}  
const double circleSurface = getCircleSurface(0.5);
```

or

```
const float oneEighty = degreesToRadians(180.0f);
```

Since C++14, constexpr functions may have more than one line.

Call other constructors

In C++03, other constructors could not be called. Workaround: call a common member function

```
class Widget {  
    int number_;  
    void construct(int number) { number_ = number; }  
public:  
    Widget(int number) { construct(number); }  
    Widget() { construct(42); }  
};
```

Call other constructors

Wrong workaround:

```
class Widget {
    int number_;
public:
    Widget(int number) : number_(number) { }
    Widget() {
        Widget(42);
        // this->number_ = 0 or something undefined
    }
};
```

What is wanted but does not compile:

```
Widget() : Widget(42) { }
```

Call other constructors

11

In C++11 other peer constructors may be called (“delegation”).

```
class Widget {  
    int number_  
public:  
    Widget(int number) : number_(number) { }  
    Widget() : Widget(42) { }  
};
```

An object is constructed once *any* constructor finishes execution. Since multiple constructors are allowed to execute, each delegating constructor will be executing on a fully constructed object of its own type.

Call or import base class constructors

11

Call base class constructor:

```
struct Base {  
    Base(int number);  
};  
struct Derived : Base {  
    Derived(int number) : Base(number) { }  
};
```

Import all base class constructors:

```
struct Derived : Base {  
    using Base::Base  
};  
// can call Derived(42);
```

Override a base class method

```
struct Animal {
    char* say() const { return "??"; }
};

struct Cow : Animal {
    char* say() const { return "moo"; }
};

struct Pig : Animal {
    char* say() const { return "oink"; }
};

std::vector<Animal*> animals{new Cow, new Pig};
for(const Animal* a : animals) {
    std::cout << a->say() << std::endl;
}
```

Prints "??" for both the cow and the pig.

Dynamic polymorphism

```
struct Animal {  
    virtual char* say() const { return "??"; }  
};  
struct Cow : Animal {  
    virtual char* say() const { return "moo"; }  
};  
struct Pig : Animal {  
    virtual char* say() const { return "oink"; }  
};
```

Prints “moo” and “oink”. What if `const` is forgotten in the Pig’s `say()`? It will compile and print “moo” and “??”.

Typo in derived class

The following code compiles but contains no virtual function overrides.

```
struct Base {  
    void f1() const;  
    virtual void f2() const;  
    virtual void f3(int x);  
}  
  
struct Derived : Base {  
    void f1() const;  
    virtual void f2();  
    virtual void f3(unsigned int x);  
}
```

overrides

11

Make explicit that a derived class is expected to override a base class:

```
struct Derived : Base {  
    void f1() const override;  
    virtual void f2() override;  
    virtual void f3(unsigned int x) override;  
}
```

The compiler will complain about all the overriding-related issues.

final

11

Prevent overriding:

```
struct Base {  
    virtual void f() final;  
};  
struct Derived : Base {  
    virtual void f(); // error: 'f' overrides a 'final' function  
};
```

Prevent derivation:

```
struct Base final { };  
struct Derived : Base { }; // error: 'Base' is marked 'final'
```

This could be achieved in C++03 with private virtual inheritance.

Explicitly deleted functions

11

```
struct Widget {  
    void f(double i);  
    void f(int) = delete;  
};  
  
Widget w;  
w.f(3.0); // ok  
w.f(3);   // error: explicitly deleted
```

Explicitly deleted special member functions:

```
struct Widget {  
    Widget() = default;  
    Widget(const Widget&) = delete;  
    Widget& operator=(const Widget&) = delete;  
};
```

Unrestricted unions



In C++03, unions cannot contain any objects that define a non-trivial constructor or destructor. C++11 lifts some restrictions.

```
struct Widget {  
    int x_;  
    Widget() {}  
    Widget(int x) : x_(x) {}  
};
```

```
union U {  
    double f;  
    Widget w; // illegal in C++03, legal in C++11  
};
```

```
U u; // error: call to implicitly deleted default constructor  
// note: default constructor of 'U' is implicitly deleted  
// because field 'w' has a non-trivial default constructor
```

Unrestricted unions



A constructor for the union must be manually defined:

```
union U {  
    double f;  
    Widget w;  
    U(int x) : w(x) { }  
};  
U u(3);
```

Strongly typed enumerations

11

C++03 enumerations are not type-safe. It allows the comparison between two enum values of different enumeration types. Type-safe enumeration:

```
enum class Enum1 {  
    Val1 = 100,  
    Val2 // = 101  
};
```

Override the default underlying type:

```
enum class Enum2 : unsigned long {Val1, Val2};
```

Other examples:

```
enum Enum3; // Underlying type cannot be determined (C++03/11)  
enum class Enum4; // Underlying type is int (C++11)
```

Template type deduction for classes

Template class:

```
template<typename T> // "typename" and "class" are synonymous
class Stack {
    std::vector<T> elems_;
public:
    void push(const T& elem) { elems_.push_back(elem); }
    T top() const { return elems_.back(); }
    void pop() { elems_.pop_back(); }
};
```

```
Stack<double> myStack; // T = double
```

```
Stack<int*> myStack; // T = int*
```

One Stack code for double and another Stack code for int*.

Template type deduction for functions

```
template <typename T>
inline T max(T a, T b) {
    return a > b ? a : b;
}

std::cout << max(3.0, 7.2) << std::endl; // T = double
std::cout << max("hello", "world") << std::endl;
```

```
struct Widget {
    int x_;
    Widget(int x) : x_(x) { }
};

Widget w1(3), w2(4);
std::cout << max(w1, w2) << std::endl;
// error: invalid operands to binary expression
```

Template partial specialization for loop unrolling

R

The template may also be a value instead of a class, allowing e.g. for compile-time loop unrolling, with partial specialization:

```
template <unsigned int N>
int factorial() {
    return N * factorial<N - 1>();
}
```

```
template <>
int factorial<0>() {
    return 1;
}

std::cout << factorial<4>() << std::endl; // yields 24
```

Since C++11, the same can be performed with constexpr.

auto for variable type deduction

11

```
std::vector<int> vec{10, 11, 12, 13};  
for(auto it = vec.cbegin(); it != vec.cend(); ++it) {  
    std::cout << *it << std::endl;  
}
```

decltype for return type deduction

11

```
template<typename T, typename U>
SomeType mul(T x, U y) {
    return x*y;
}
```

The return type is “the type of $x*y$ ”. How can we write that?

```
template<typename T, typename U>
decltype(x*y) mul(T x, U y) {
    return x*y;
}
```

It does not compile because x and y are not in scope. We can use a trailing return type.

decltype for return type deduction

11

Trailing return type:

```
template<typename T, typename U>
auto mul(T x, U y) -> decltype(x*y) {
    return x*y;
}
```

When `decltype` is read, `x` and `y` are now known to the compiler. `auto` is only meant to tell the compiler that it will find the return type after `->`. `auto` does not deduce anything.

Generalized return type deduction

14

Since C++14, `auto` can deduce the return type of functions:

```
auto floor(double x) { return static_cast<int>(x); }
```

Multiple returns are allowed, but the type must be the same:

```
auto f() {  
    while(something()) {  
        if(expr) {  
            return foo() * 42;  
        }  
    }  
    return bar.baz();  
}
```

Generalized return type deduction

14

The return type deduction also works for recursive functions, provided the non-recursive return statement is before the recursive call.

```
auto f() { return f(); } // error: return type of f is unknown
```

```
auto sum(int i) {  
    if (i == 1)  
        return i;           // return type deduced to int  
    else  
        return sum(i-1)+i; // ok to call it now  
}
```

decltype(auto)

14

auto always deduces a non-reference type (int, double...). auto& always deduces a reference type (int&, double&...). But auto cannot conditionnaly see if the return type is a reference.

decltype can, but requires to write an expression in its parentheses.

We can use the best of both worlds with decltype(auto) which can deduce if the type is a reference.

```
template<typename T, typename U>
decltype(auto) mul(T x, U y) {
    return x*y;
}
```

Copy-and-swap idiom

```
class Array {
    int* arr_;
    unsigned int size_;
public:
    Array(unsigned int size = 0) : size_(size),
        arr_(size > 0 ? new int[size]() : nullptr) { }

    Array(const Array& other) : size_(other.size_),
        arr_(other.size_ > 0 ? new int[other.size_] : nullptr)
    {
        std::copy(other.arr_, other.arr_ + size_, arr_);
    }

    ~Array() { delete[] arr_; }
};
```

Copy-and-swap idiom

Naïve implementation of operator=:

```
Array& operator=(const Array& other) {  
    if (this != &other) {  
        delete [] arr_;  
        arr_ = nullptr;  
  
        arrSize_ = other.arrSize_;  
        arr_ = arrSize_ ? new int[arrSize_] : nullptr;  
        std::copy(other.arr_, other.arr_ + arrSize_, arr_);  
    }  
    return *this;  
}
```

Copy-and-swap idiom

Issues:

- ❶ code duplication;
- ❷ `if(this != &other)` mandatory but should not occur;
- ❸ if `new[]` fails, `*this` will have been modified (no strong exception guarantee).

Copy-and-swap idiom

Successful solution:

```
friend void swap(Array& first, Array& second) {
    using std::swap; // enable ADL
    swap(first.size_, second.size_);
    swap(first.arr_, second.arr_);
}

Array& operator=(Array other) {
    swap( *this, other );
    return *this;
}
```

```
Array getArray(unsigned int size) { return Array(size); }
```

```
Array arr = getArray(4);
```

Copy Elisions, Returned Value Optimization

Distinguish *parameter* and *argument*:

- other is the *parameter* of the function;
- the array returned by `getArray(4)` is the *argument*, from which the parameter `other` is instantiated.

If the argument is a temporary object which will be destroyed, why make a copy for `other`? The argument should *become* `other`, thereby avoiding copy.

Most C++03 compilers could grasp this optimization opportunity whenever possible and elude the copy.

Guideline: Do not copy the function arguments. Instead, pass them by value and let the compiler manage the copying.

Lvalue – Rvalue

11

C++11 formalized everything:

- a temporary object, eligible for copy elision, is an *rvalue*;
- otherwise it is an *lvalue*.

References:

- an integer *lvalue-reference* is denoted `int&`;
- an integer *rvalue-reference* is denoted `int&&`.

Move constructor and move assignment operator

11

There are two new special member functions.

Move constructor:

```
Array(Array&& other) : Array() {  
    swap( *this, other );  
}
```

Move assignment operator:

```
Array& operator=(Array&& other) {  
    swap( *this, other );  
    return *this;  
}
```

Specialize code for lvalue or rvalue

11

```
void f(int& param) {
    std::cout << "lvalue" << std::endl;
}

void f(int&& param) {
    std::cout << "rvalue" << std::endl;
}

int a;
f(a); // calls void f(int&)
f(std::move(a)); // calls void f(int&&)
```

If you can take a variable's address, it usually is an lvalue. Otherwise it is usually an rvalue.

std::move

11

`std::move` does not move anything. It casts its parameter as an rvalue.

```
struct Widget {  
    Array arr_;  
    Widget(const Array& param) : arr_(std::move(param)) { }  
};  
Array arr(3);  
Widget w(arr);
```

compiles but calls the copy constructor of `Array` instead of the move constructor. `std::move` returns a `const Array&&` which cannot be casted as a non-`const Array&&`, but can be casted as a `const Array&`.

```
Array(const Array& other); // copy constructor  
Array(Array&& other);      // move constructor
```

Distinguish forwarding from rvalue references

11

A *forwarding reference* is either an lvalue reference or an rvalue reference.
A forwarding reference has the same syntax as an rvalue reference, namely “&&”.

Rvalue references examples:

```
void f(Widget&& param);
```

```
Widget&& w1 = Widget();
```

Forwarding references:

```
template<typename T>  
void f(T&& param);
```

```
auto&& w2 = w1;
```

Distinguish forwarding from rvalue references

11

Both forwarding references have in common the presence of type deduction. Rvalue references do not.

```
template<typename T>
void f(T&& param) {
    // Here we can take param's address
    // so param is always an lvalue
}

Widget w;
f(w); // param's type is lvalue-reference Widget&
      // param is lvalue
      // T = Widget&

f(std::move(w)); // param's type is rvalue-reference Widget&&
                 // param is lvalue
                 // T = Widget
```

std::forward

11

```
void process(Widget& lValArg); // for lvalues
```

```
void process(Widget&& rValArg); // for rvalues
```

```
template<typename T>
```

```
void logAndProcess(T&& param) {
```

```
    auto now = std::chrono::system_clock::now();
```

```
    makeLogEntry("Calling 'process'", now);
```

```
    process(param); // always calls process(Widget&)
```

```
}
```

```
Widget w;
```

```
logAndProcess(w);
```

```
logAndProcess(std::move(w));
```

Because param is always an lvalue, process(Widget&) is always called.

std::forward

11

If param's type is an rvalue-reference, we need to cast the lvalue param as an rvalue. This is what `std::forward` does.

`std::move` always casts as an rvalue; `std::forward` conditionally casts as an rvalue.

But how to distinguish? Look at T:

- 1 if param is an lvalue-reference, `T = Widget&`;
- 2 if param is an rvalue-reference, `T = Widget`.

T must be given to `std::forward` so it can decide whether to cast its parameter as an rvalue or not.

```
process(std::forward<T>(param));
```

Subtle forwarding and rvalue references

11

R

```
template<typename T>
class vector {
public:
    void push_back(T&& x);
};
```

is not a forwarding reference. Writing `std::vector<Widget> v;` causes the template to be instantiated as

```
class vector<Widget> {
public:
    void push_back(Widget&& x);
};
```

which involves no type deduction.

Subtle forwarding and rvalue references



A template forwarding reference must have the form `T&&`. Hence, the two following references do not qualify for being forwarding references.

```
template<typename T>  
void f(const T&& param);
```

```
template<typename T>  
void f(std::vector<T>&& param);
```

Subtle forwarding and rvalue references



The variadic templates of `emplace` involve type deduction.

```
template<typename T>
class vector {
public:
    template<typename... Args>
    void emplace(Args&& args);
};
```

Implementation of std::move in C++11



```
template<typename T>
typename remove_reference<T>::type&& move(T&& param) {
    using ReturnType = typename remove_reference<T>::type&&;
    return static_cast<ReturnType>(param);
}
```

param is a forwarding reference:

- if param is an lvalue-reference, $T = \text{Widget\&}$;
- if param is an rvalue-reference, $T = \text{Widget}$.

Steps:

- 1 remove_reference<T>::type returns Widget;
- 2 remove_reference<T>::type&& returns Widget&&;
- 3 static_cast<Widget&&> casts param as an rvalue.

Implementation of std::move in C++14

14

R

std::remove_reference_t<T> replaces
typename remove_reference<T>::type.

std::move can elegantly be written as

```
template<typename T>
decltype(auto) move(T&& param) {
    using ReturnType = remove_reference_t<T>&&;
    return static_cast<ReturnType>(param);
}
```

Lambda expressions

11

```
auto lambda = [](int x, int y) -> int { return x + y; };  
std::cout << lambda(2, 3) << std::endl;
```

Components:

- (int x, int y) is the parameters list;
- int after -> is the return type;
- {...} is the lambda body.

A lambda can be executed upon declaration with trailing "()":

```
std::cout  
  << [](int x, int y) -> int { return x + y; }(2, 3)  
  << std::endl;
```

std::for_each and std::transform

11

Apply a lambda to each element of an iterator:

```
void f(std::vector<int>& v) {  
    std::for_each(v.begin(), v.end(),  
        [](int p) { std::cout << p << std::endl; });  
}
```

Modify each element with a lambda:

```
void f(std::vector<double>& v) {  
    std::transform(v.begin(), v.end(), v.begin(),  
        [](double d) { return d < 0.00001 ? 0 : d; });  
}
```

C++03 already had `std::for_each` and `std::transform`, but cumbersome functors were necessary since lambdas were not available.

std::for_each and std::transform

11

If the body is not a single return expression, the type has to be explicitly stated:

```
void f(std::vector<double>& v) {  
    std::transform(v.begin(), v.end(), v.begin(),  
        [](double d) -> double {  
            if (d < 0.0001)  
                return 0;  
            else  
                return d;  
        });  
}
```

Lambda captures

11

[] allows to capture other variables from the scope.

```
void f(std::vector<double>& v, double epsilon) {
    std::transform(v.begin(), v.end(), v.begin(),
        [epsilon](double d) -> double {
            if (d < epsilon)
                return 0;
            else
                return d;
        });
}
```

We can capture by both reference and value.

Generalized lambda captures

14

An element of the capture can be initialized with `=`. This allows renaming of variables and to capture by moving.

```
int x = 4;
auto y = [&r = x, x = x+1]() -> int {
    r += 2;
    return x+2;
}(); // Updates ::x to 6, and initializes y to 7
```

```
std::unique_ptr<int> ptr(new int(10));
auto lambda = [value = std::move(ptr)] { return *value; };
```

`std::unique_ptr` can be moved but not copied. Capture by move is the only way to capture a `std::unique_ptr`.

Generic lambdas

14

Lambda function parameters may be declared auto.

```
auto lambda = [](auto x, auto y) { return x + y; };
```

Improved return type deduction

14

C++14 allows deduced return types for every function, not only those of the form `return expression`.

```
void f(std::vector<double>& v, double epsilon) {  
    std::transform(v.begin(), v.end(), v.begin(),  
        [epsilon](auto d) {  
            if (d < epsilon)  
                return 0;  
            else  
                return d;  
        });  
}
```

Remaining modern C++ topics

- type traits (`std::enable_if`, `std::is_pointer...`);
- smart pointers;
- threading facilities;
- `noexcept`;
- `std::function`;
- `std::bind` and `std::ref`;
- `std::regex`;
- extern template;
- inline namespace...